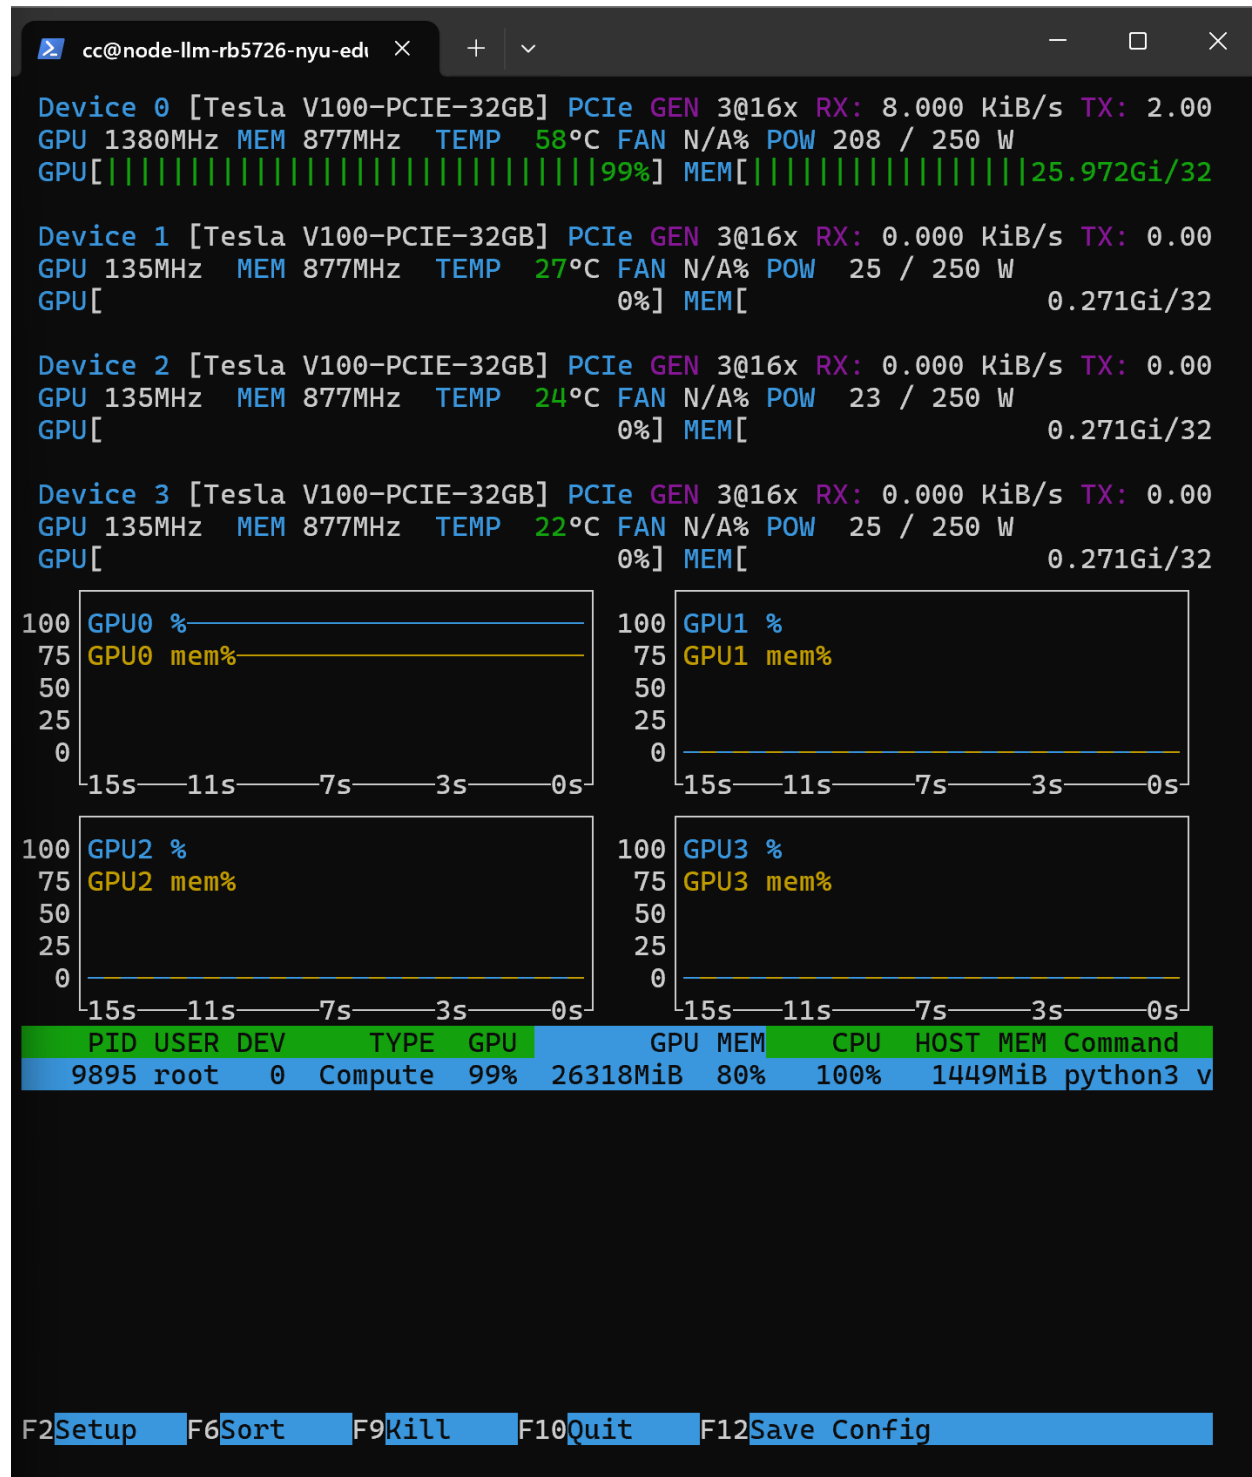


## TinyLlama 1.1B model on a single V100 32GB



nvidia-smi

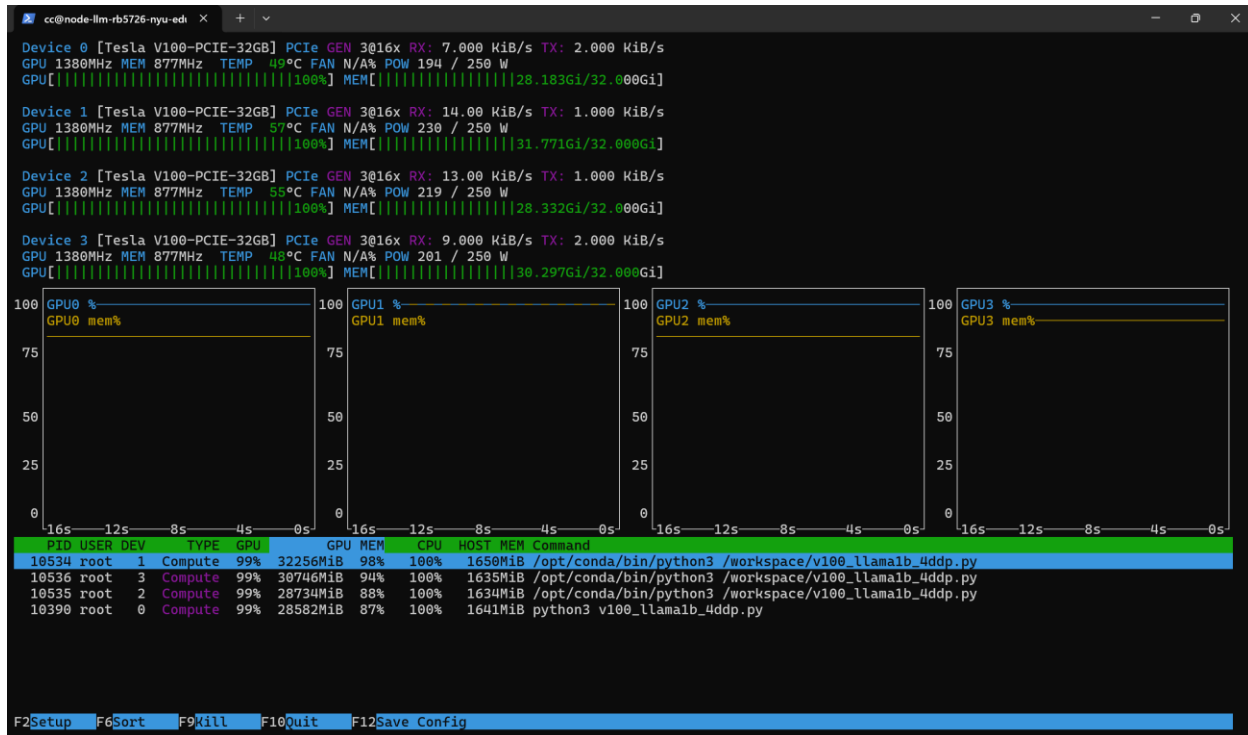
```

root@fd12a393fef3: /workspa × + ∨
Epoch 0: 85%| 106/125 [04:22<00:47, 0.40it/s, v_num=1, train_loss=1.
Allocated: 16.00 GB, Reserved: 25.33 GB
Epoch 0: 86%| 107/125 [04:24<00:44, 0.41it/s, v_num=1, train_loss=1.
Allocated: 19.49 GB, Reserved: 25.33 GB
Epoch 0: 86%| 108/125 [04:26<00:41, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.92 GB, Reserved: 25.33 GB
Epoch 0: 87%| 109/125 [04:29<00:39, 0.40it/s, v_num=1, train_loss=1.
Allocated: 16.68 GB, Reserved: 25.33 GB
Epoch 0: 88%| 110/125 [04:30<00:36, 0.41it/s, v_num=1, train_loss=1.
Allocated: 19.62 GB, Reserved: 25.33 GB
Epoch 0: 89%| 111/125 [04:33<00:34, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.91 GB, Reserved: 25.33 GB
Epoch 0: 90%| 112/125 [04:36<00:32, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.91 GB, Reserved: 25.33 GB
Epoch 0: 90%| 113/125 [04:38<00:29, 0.41it/s, v_num=1, train_loss=1.
Allocated: 19.55 GB, Reserved: 25.33 GB
Epoch 0: 91%| 114/125 [04:41<00:27, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.91 GB, Reserved: 25.33 GB
Epoch 0: 92%| 115/125 [04:44<00:24, 0.40it/s, v_num=1, train_loss=0.
Allocated: 23.91 GB, Reserved: 25.33 GB
Epoch 0: 93%| 116/125 [04:47<00:22, 0.40it/s, v_num=1, train_loss=1.
Allocated: 19.58 GB, Reserved: 25.33 GB
Epoch 0: 94%| 117/125 [04:49<00:19, 0.40it/s, v_num=1, train_loss=1.
Allocated: 21.80 GB, Reserved: 25.33 GB
Epoch 0: 94%| 118/125 [04:51<00:17, 0.40it/s, v_num=1, train_loss=1.
Allocated: 23.93 GB, Reserved: 25.33 GB
Epoch 0: 95%| 119/125 [04:54<00:14, 0.40it/s, v_num=1, train_loss=1.
Allocated: 23.93 GB, Reserved: 25.33 GB
Epoch 0: 96%| 120/125 [04:57<00:12, 0.40it/s, v_num=1, train_loss=1.
Allocated: 12.95 GB, Reserved: 25.33 GB
Epoch 0: 97%| 121/125 [04:58<00:09, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.42 GB, Reserved: 25.33 GB
Epoch 0: 98%| 122/125 [05:01<00:07, 0.40it/s, v_num=1, train_loss=1.
Allocated: 15.96 GB, Reserved: 25.33 GB
Epoch 0: 98%| 123/125 [05:03<00:04, 0.41it/s, v_num=1, train_loss=1.
Allocated: 23.92 GB, Reserved: 25.33 GB
Epoch 0: 99%| 124/125 [05:06<00:02, 0.41it/s, v_num=1, train_loss=1.
Allocated: 19.92 GB, Reserved: 25.33 GB
Epoch 0: 100%| 125/125 [05:08<00:00, 0.41it/s, v_num=1, train_loss=1.
`Trainer.fit` stopped: `max_epochs=1` reached.
Epoch 0: 100%| 125/125 [05:08<00:00, 0.41it/s, v_num=1, train_loss=1.
root@fd12a393fef3: /workspace#


```

Training time

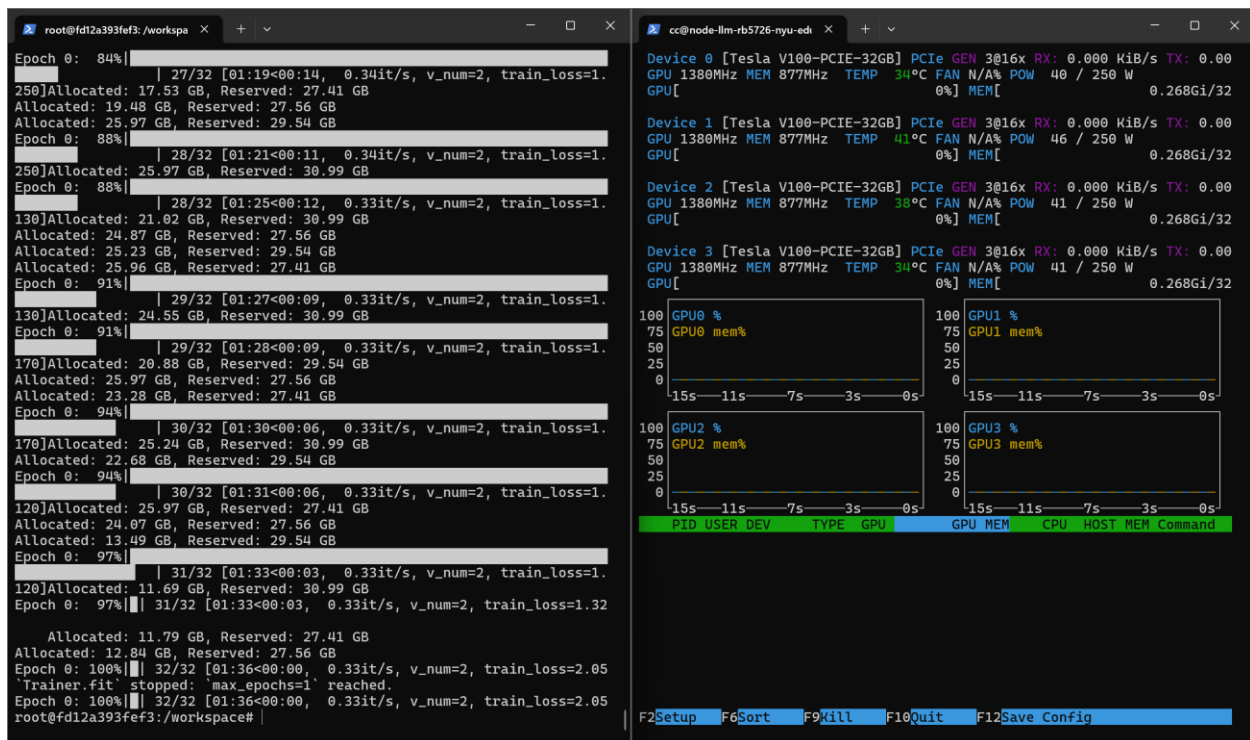
## TinyLlama 1.1B model on 4x V100 32GB with DDP



nvidia-smi



Training time



After completion of TinyLlama 1.1B model on 4x V100 32GB with DDP